

AI Notes

Context

Exam :

Monday 30th , 9:30am – 12:00pm

4 Questions :

1. Clustering
2. Supervised Learning
3. Search
4. Optimisation

Each 15 marks,
total 60 marks

****note: no GMM and weka**

Content :

- Clustering (Week 1 & 2)
- Supervised Learning (Week 3 - 6)
- Search (Week 7 & 8)
- Optimisation (Week 9 & 10)

Clustering :

- Hierarchical Clustering
- K-means
- GMM/EM
- DBSCAN

Supervised Learning :

- Linear Regression
- Gradient Descent
- Logistic Regression
- Neural Network
- Evaluate & Hyperparameter Tuning
- Naive Bayes
- KNN-algorithm

Search :

- Search Problem
- BFS
- DFS
- A*

Optimisation :

- Hill Climbing
- Simulated Annealing
- Formulation
- Constraint handling

Introduction

Basic Introduction:

Assuming a data set with n samples, and each sample has d dimensions :

$$[(x_1^1, x_2^1, \dots, x_d^1), (x_1^2, x_2^2, \dots, x_d^2), \dots, (x_1^n, x_2^n, \dots, x_d^n)] \\ = 1, 2, \dots, n$$

Instance : One data sample

Dimension : Number of coordinates to specify the samples

Feature / Attribute : The value at each dimension

you can think of n samples as data inputs, and d dimensions as attributes
(or in Objects terms,
 n : Objects
 d : properties)

There are 3 types of Machine Learning:

- **Supervised Learning**

(Labelled Data)

- **Unsupervised Learning**

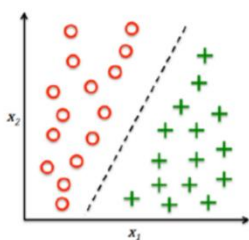
(Unlabeled Data, may be due to data being too big / unclear)

- **Reinforced Learning**

(Learns from rewards, rewards changes by time)

Classification

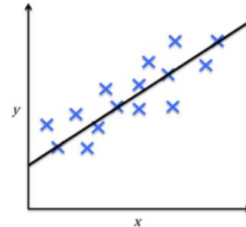
Outputs are **categorical**



eg: handwritten numbers $\rightarrow [1, 2, 3, \dots]$

Regression

Outputs are **continuous numbers**



eg: Student Score Marks $\rightarrow 83.2\%$

Clustering

Hierarchical Clustering:

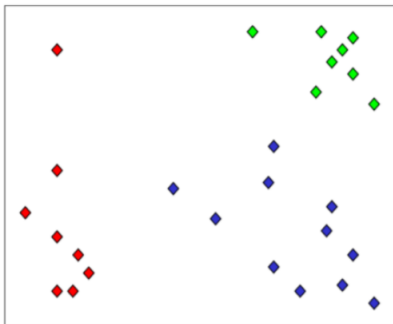
Terminology

Singleton : one data point

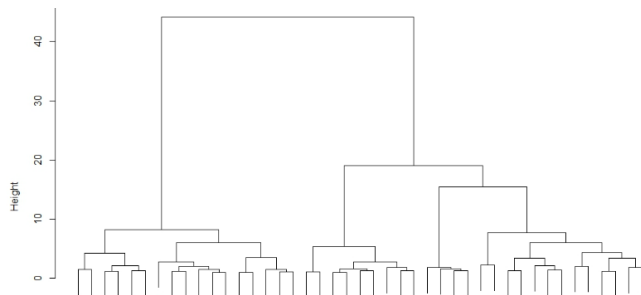
Cluster : multiple data points

Segments data into clusters such that there is a “natural grouping” among objects, it creates a **hierarchical decomposition** of the set of objects using some criterions

eg :



And produces a **dendrogram** :



Note : the distance between clusters are noted as height

In the end, there will be **only 1 cluster** all grouped together

Measuring Distance

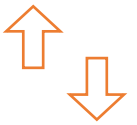
we will be using **Euclidean Distance** only

$$\text{distance}(a, b) = \sqrt{(a_x - b_x)^2 + (a_y - b_y)^2}$$

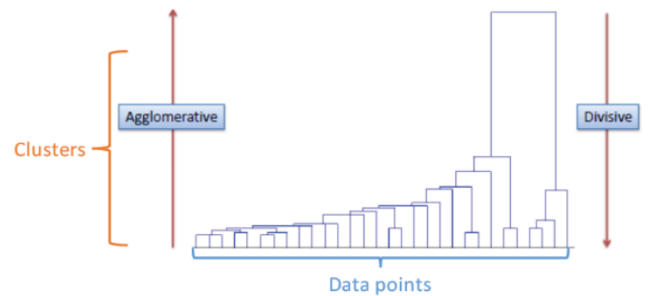
extra : it is also Minkowski with dimension = 2

There are 2 ways of Hierarchical clustering:

- **Agglomerative** (*bottom-up merging*)
- **Divisive** (*top-down merging*)

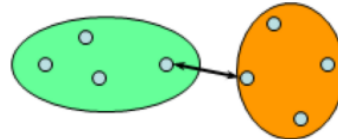


Divisive is not recommended as it can be computationally expensive



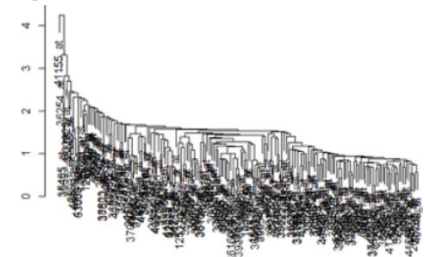
There are 3 ways to measure distance between clusters:

- **Single Linkage**



- Distance of the **closest pair**
- Produces **long chain shape** of clusters

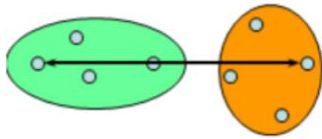
Eg:



$$c_1, c_2 = \min(\forall \text{distance}(p_{c_1}, p_{c_2}))$$

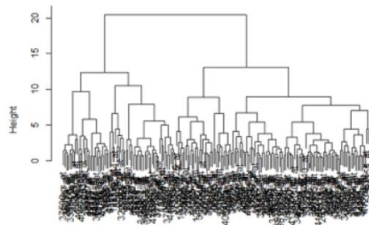
Where c_1, c_2 are clusters and return **minimum** distance between 2 points in c_1 and c_2

Complete Linkage



- Distance of the **furthest pair**
- Produces **compact shape** clusters
- **Sensitive to noise**

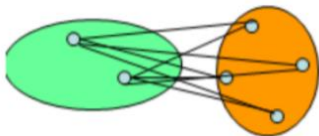
Eg:



$$c_1, c_2 = \max(\forall \text{distance}(p_{c_1}, p_{c_2}))$$

return **maximum** distance between 2 points in c_1 and c_2

Group Average



- **Average of all distances**
- Most used
- **Robust against noise**

$$c_1, c_2 = \frac{\sum \forall \text{distance}(p_{c_1}, p_{c_2})}{\text{Num } p_{c_1} + \text{Num } p_{c_2}}$$

return **Average** distance of **all the points** in c_1 and c_2

Advantages

- Deterministic results (same result for any run)
- Does not need to specify number of clusters
- Can create cluster of arbitrary shapes

Disadvantages

- Does not scale up to larger datasets, **time complexity is $O(n^2)$**
- Will impose hierarchical structure even though data is not appropriate for it
- Once a decision is made, can't be undone

Algorithm (agglomerative):

Step 1 : Find every Euclidean distance

(create distance matrix)

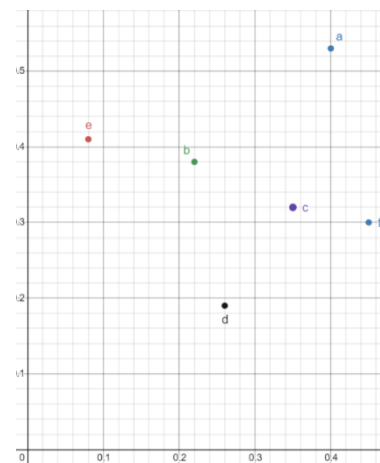
Step 2 : Get the **minimum distance**, which the 2 points/cluster joins becoming the new cluster

Step 3 : Recalculate the distance matrix with the new cluster

- Complete Linkage → **Max**
- Single Linkage → **Min**
- Group Average → **Average**

Step 4 : Repeat **Step 2**, until only 1 cluster left, **Done!**

Example of a run (Complete Linkage):



point	x-axis	y-axis
a	0.40	0.53
b	0.22	0.38
c	0.35	0.32
d	0.26	0.19
e	0.08	0.41
f	0.45	0.30

Step 1 – Use Euclidean distance to create distance matrix

(in this case we are only calculating point a and b)

$$\text{distance}(a, b) = \sqrt{(a_x - b_x)^2 + (a_y - b_y)^2}$$

$$\text{distance}(a, b) = \sqrt{(0.40 - 0.22)^2 + (0.53 - 0.38)^2}$$

$$= \sqrt{(0.18)^2 + (0.15)^2}$$

$$= \sqrt{0.0549}$$

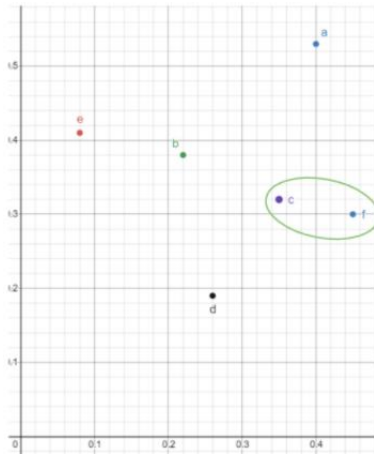
$$= 0.23$$

Distance Matrix:

	a	b	c	d	e	f
a	0					
b	0.23	0				
c	0.22	0.15	0			
d	0.37	0.20	0.15	0		
e	0.34	0.14	0.28	0.29	0	
f	0.23	0.25	0.11	0.22	0.39	0

Step 2 – get minimum value, forms the cluster

	a	b	c	d	e	f
a	0					
b	0.23	0				
c	0.22	0.15	0			
d	0.37	0.20	0.15	0		
e	0.34	0.14	0.28	0.29	0	
f	0.23	0.25	0.11	0.22	0.39	0



Step 3 – Recalculate the distance matrix (Complete Link)

to update the table:

$$\text{New Value} = \max(\text{distance}(p_1, n), \text{distance}(p_2, n))$$

Which concludes:

$$\max(\text{distance}(c, a), \text{distance}(f, a)) = 0.23$$

$$\max(\text{distance}(c, b), \text{distance}(f, b)) = 0.25$$

$$\max(\text{distance}(c, d), \text{distance}(f, d)) = 0.22$$

$$\max(\text{distance}(c, e), \text{distance}(f, e)) = 0.39$$

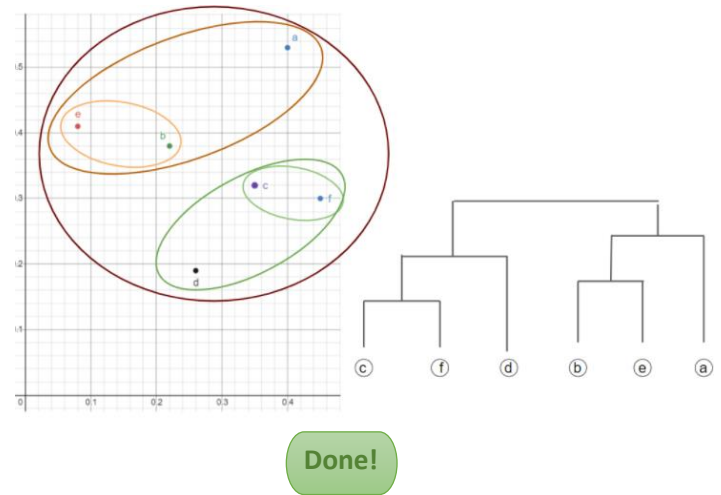
Updated Table :

	a	b	c,f	d	e
a	0				
b	0.23	0			
c,f	0.23	0.25	0		
d	0.37	0.20	0.15	0	
e	0.34	0.14	0.28	0.29	0

Step 4 – Repeat Step 2 until 1 cluster left

Will give an output of the distance of the final 2 cluster

**many more steps here but will skip ahead till 1 cluster left as it's just a repeat



K-Means:

Terminology

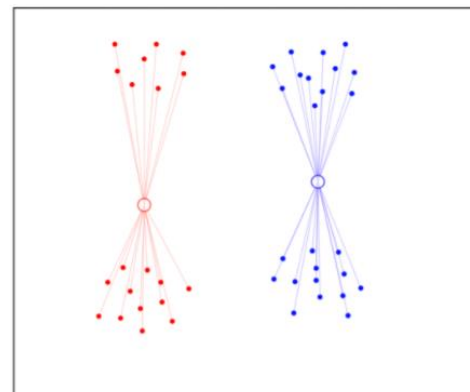
K : number of clusters

Centroid : the center of uniform density

Centroid Based – Describes each cluster by its **mean**

Its objective is to *minimize the within cluster variances* of all clusters that you set, **K**

eg (**K** = 2) :



Advantages

- Easy to implement
- Efficient

Disadvantages

- Non-deterministic results (different result for every run)
- May result in **local optima** (multiple restart to get global optima)
- **Require number of clusters** in advance

Algorithm :

Step 1 : Set number of cluster , **K**

(K will usually have coordinates, if not → set random)

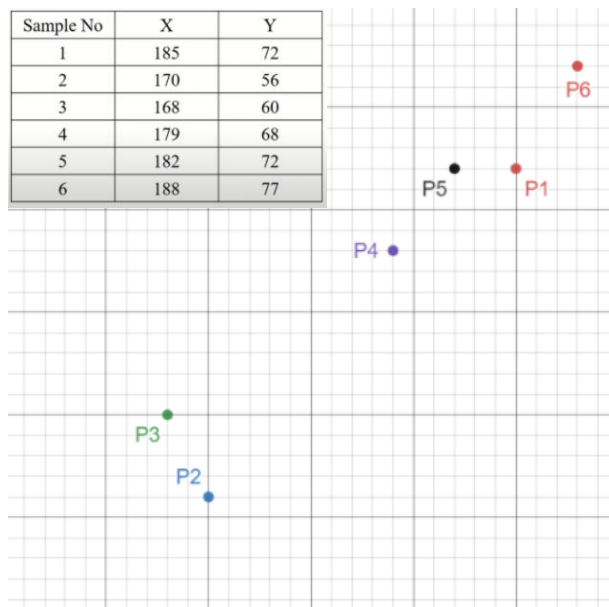
Step 2 : Calculate the distance matrix of every points to **K centroid**

Step 3 : Assign the points to the closest K centroid (minimum)

Step 4 : Recalculate the cluster centroid mean

Step 5 : Repeat **Step 2**, until **K-mean** stop changing, Done!

Example of a run :



Step 1 : Set number of cluster , **K**

Set **K = 2** and in this case: $K_1 = P_1, K_2 = P_2$

Initial Centroid		
Cluster	X	Y
k1	185	72
k2	170	56

Step 2 : Calculate the distance matrix of every points to **K centroid**

Using Euclidean Distance formula :

$$\text{distance}(a, b) = \sqrt{(a_x - b_x)^2 + (a_y - b_y)^2}$$

Gives :

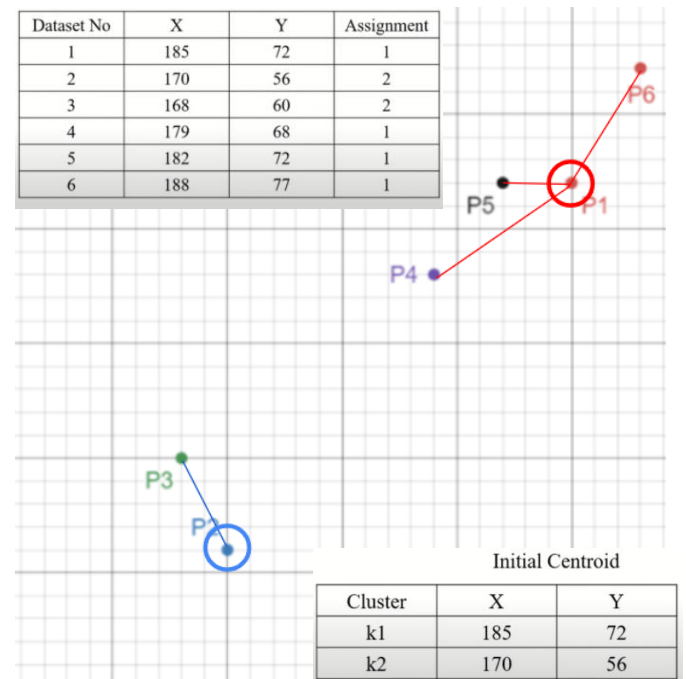
Sample No	K1	K2
3	20.809	4.472
4	7.211	15
5	3	20
6	5.831	27.659

Step 3 : Assign the points to the closest K centroid

Using formula :

$$\min(\text{distance}(P_n, K_1), \text{distance}(P_n, K_2))$$

"assign to K closest to it"



Step 4 : Recalculate the cluster centroid mean

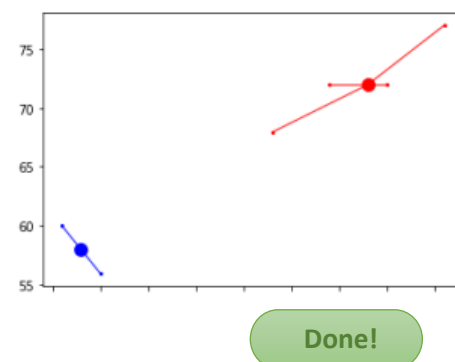
Formula for new K-mean if dimension = 2 :

$$\frac{\sum P_x}{n}, \frac{\sum P_y}{n}$$

n = number of points in that cluster

Cluster	X	Y
k1	$= (185+179+182+188) / 4$ $= 183$	$= (72+68+72+77) / 4$ $= 72$
k2	$= (170 + 168) / 2$ $= 169$	$= (60 + 59) / 2$ $= 58$

Step 5 : Repeat **Step 2**, until **K-mean** stop changing



GMM/EM:

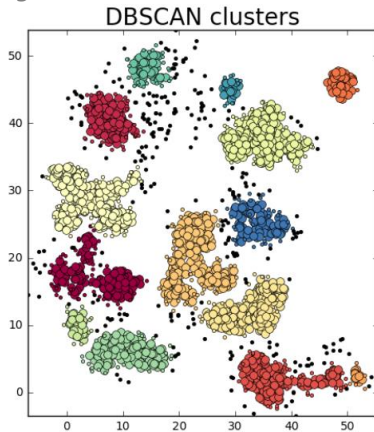
NO

DBSCAN:

Density Based Spatial Clustering of Application with Noise (you do not need to memorize the name)

Discover clusters by density of regions

Eg:



Usually, there are 2 parameters :

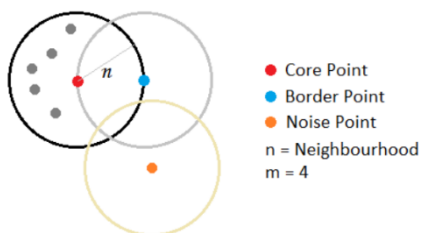
- Radius of circle (ϵ)
- Minimum number of points in that circle ($minPts$)

Points are categorized into 3 terms:

Core : point has $minPts$ and within radius ϵ

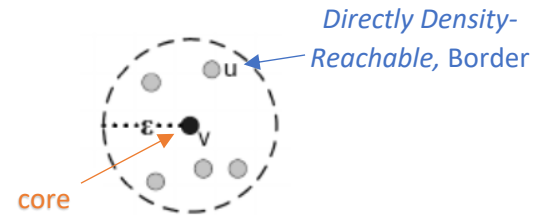
Border : does not have $minPts$ and point lies within radius ϵ of core

Noise/Outlier : point which are not border or core



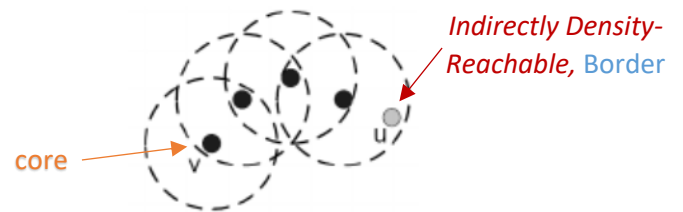
Forming Clusters

Directly Density-Reachable



A point, u is *Directly Density-Reachable* if it is within a core's radius of point v

Indirectly Density-Reachable



A point, u is *Indirectly Density-Reachable* if it is a border but not within a core's radius of point v

Advantages

- Can discover arbitrary shape clusters
- Robust to noise/outliers
- Doesn't need number of cluster

Disadvantages

- Non-deterministic results (border points may have different clusters depending on code)
- If data is not well-explained, choosing ϵ and $minPts$ may be difficult
- May fail if data too sparse

Algorithm :

Step 1 : Label all the points , **core/border/noise**

Step 2 : Remove all the noise/outlier points

(or set them as Noise)

Step 3 : For every core point, Assign points that are *density-Reachable* from that core and make it a cluster

Step 4 : Reassign border points belonging to more than 1 cluster to be the one closest to it

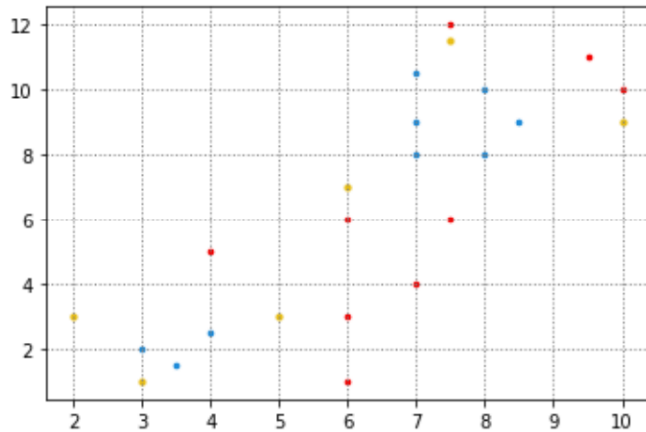
Done!

Example of a run :

Step 1 : Label all the points , **core**/**border**/**noise**

have set:

- Radius of Circle (ϵ) $\rightarrow 1.5$
- Number of Neighbors (minPts) $\rightarrow 4$

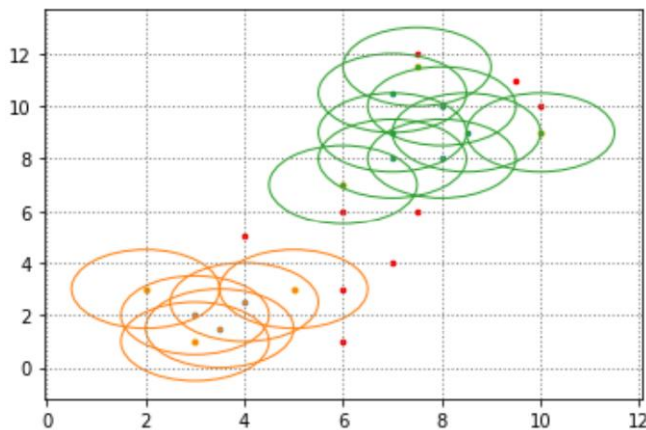


Step 2 : Remove all the noise/outlier points

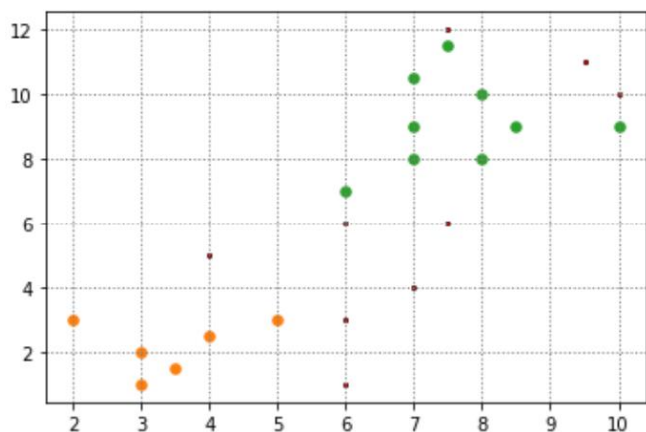
(in code, we can just ignore them to increase efficiency of program)

```
onlyList = core + border
```

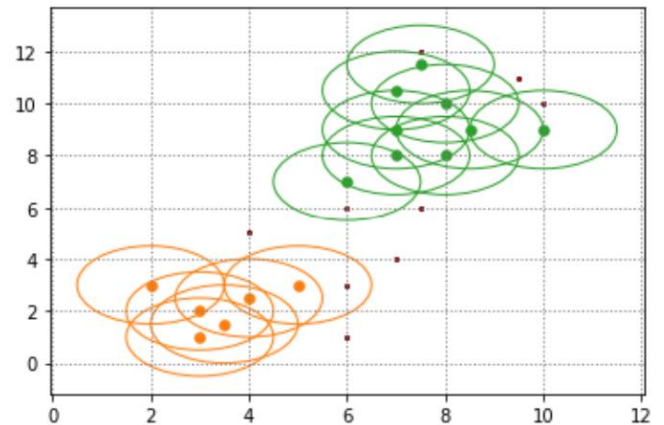
Step 3 : For every core point, Assign points that are density-Reachable from that core and make it a cluster



Without radius shown :



Step 4 : Reassign border points belonging to more than 1 cluster to be the one closest to it



(In this case, we don't need to as border doesn't belong to more than 1 cluster)

Done!