

ID 2069997

## Exam for 30203 and 30421 - Systems Programming in C/C++

After inserting your student ID and the module name in the title, header and footer, write your answers between here and the statement of good academic conduct. Your ID and the module name will automatically appear on any subsequent pages.

### Question 1

(a)

There will not be memory leakage in the program. In line 3 memory is allocated, with the address of the allocated memory stored in integer pointer A.

Although the value that A points to is modified in line 4, the memory address is not changed.

A still contains the address of the allocated memory.

Line 5 initializes the integer pointer B. A is unchanged.

In line 6 this same memory address is stored in integer pointer B by assigning B = A.

Therefore B contains the memory address of the memory allocated in line 3.

The free(int \*) function frees the allocated memory with the address that is passed to it. By free(B) in line 7, the allocated memory allocated in line 3 is freed.

There are no other allocated memories. Therefore there will not be a memory leakage.

(b)

This function returns an integer pointer pointing to unassigned values.

When this function is called, the array values are stored in stack memory and only has scope within this function. After returning the address of the start of the array (line 6), the function terminates, and the array values are deleted.

A fix could be achieved by storing the array in heap using the malloc() function.

```
int* randomArray(void)
{
    int * array = malloc(sizeof(int) * 10);
    for (int i = 0; i < 10; i++)
        array[i] = rand();
    return array;
}
```

## Question 1 (cont.)

(c)

Sum2Darray2 is able to achieve faster computation time.

In the function sum2Darray1, data is fetched with column-major order, while the function sum2Darray2 fetches data in row-major order.

When loading values from the main memory to the cache, values stored in nearby memory addresses will also be loaded into the cache to decrease the chance of the need to load values from the main memory, which is much slower.

The C compiler stores 2D arrays in row-major order.

By fetching data in row-major order, the same order as how values are stored, a higher rate of cache hit can be achieved when running sum2Darray2. This results in a higher CPU utilization rate and therefore faster computation.

## Question 2

(a)

<code>//Process P1</code>	<code>//Process P2</code>	<code>//Process P3</code>
<code>void P1(void){</code>	<code>void P2(void){</code>	<code>void P3(void){</code>
<code>wait(B);</code>	<code>Statement B</code>	<code>wait(A);</code>
<code>Statement A</code>	<code>signal (B);</code>	<code>Statement C</code>
<code>signal(A);</code>		
<code>}</code>	<code>}</code>	<code>}</code>

Trace:

```
wait(A);
wait(B);
Statement B
signal (B);
Statement A
signal(A);
Statement C
```

(b)

Race condition appears. A pointer pointing to `i` is passed to the threads, instead of a copy of the current value of `i`. Therefore, multiple threads are accessing a value in the same memory location. When a thread is created, the main thread continues to run, changes the value of `i` and creates a new thread. If this is done before the `printf()` function in the earlier thread, the pointer passed into the `printf()` function points to the updated value, and the program has an incorrect output.

## Question 2 (cont.)

(c)

```
int main(){
    pthread_t tid[4];
    int i;
    int t[4] = {0,1,2,3};

    for(i=0; i<4; i++)
        pthread create(&tid[i], NULL, foo, &t[i]);

    for(i=0; i<4; i++)
        pthread join(tid[i], NULL);

    return 0;
}
```

By assigning the values into an array, pointers pointing to different memory locations are passed into different threads. The values in the array t are also unmodified when creating new threads. This avoids race condition.

## Question 3

(a)

There is only one possible output:

2  
3

Explanation: line 6 initializes integer x = 1.

Line 7 uses the function fork() and creates a child process.

As the child process and parent process are 2 separate processes, they have their own values of x and don't interfere.

The child process, with pid == 0, executes line 8 and x = 2. The result "2" is printed on screen with a new line after it and the child process terminates.

The parent process, with pid > 0, executes line 11 and 12.

With wait(NULL) in line 11, the parent process is blocked until the child process is terminated, that is, "2" is printed on screen with a new line after it. x is assigned to be 3 and is printed out with a new line. Parent process terminates.

Question 3 (cont.)

(b)

Total turnaround time = 106

0	P2(32)	P2 arrives and gets processed
3	P2(29), P0(20)	P0 arrives
5	P2(27), P0(20), P3(3)	P3 arrives
10	P0(20), P3(3), P2(22)	Quantum time expires, P2 is forced out of CPU. P0 gets processed.
15	P0(15), P3(3), P2(22), P1(22)	P1 arrives
20	P3(3), P2(22), P1(22), P0(10)	Quantum time expires, P0 is forced out of CPU. P3 gets processed.
23	P2(22), P1(22), P0(10)	P3 gets completed, P2 gets processed.
33	P1(22), P0(10), P2(12)	Quantum time expires, P2 is forced out of CPU. P1 gets processed.
43	P0(10), P2(12), P1(12)	Quantum time expires, P1 is forced out of CPU. P0 gets processed.
49	P0(4), P2(12), P1(12), P4(29)	P4 arrives
53	P2(12), P1(12), P4(29)	P0 gets completed, P2 gets processed.
63	P1(12), P4(29), P2(2)	Quantum time expires, P2 is forced out of CPU. P1 gets processed.
73	P4(29), P2(2), P1(2)	Quantum time expires, P1 is forced out of CPU. P4 gets processed.
83	P2(2), P1(2), P4(19)	Quantum time expires, P4 is forced out of CPU. P2 gets processed.
85	P1(2), P4(19)	P2 gets completed, P1 gets processed.
87	P4(19)	P1 gets completed, P4 gets processed.
97	P4(9)	Quantum time expires, but no other processes are in queue. P4 gets processed.
106		P4 gets processed.

(c)

(i) 20495

(ii) invalid

(iii) invalid

(iv) 16385

## Question 3 (cont.)

(d)

FIFO: 8 page faults

0	1	2	3	4	2	3	0	1	4	2
0	0	0	0	4	4	4	4	4	4	4
	1	1	1	1	1	1	0	0	0	0
		2	2	2	2	2	2	1	1	1
			3	3	3	3	3	3	3	2

LRU: 9 page faults

0	1	2	3	4	2	3	0	1	4	2
0	0	0	0	4	4	4	4	1	1	1
	1	1	1	1	1	1	0	0	0	0
		2	2	2	2	2	2	2	4	4
			3	3	3	3	3	3	3	2

It may seem weird to see that LRU has more page faults than FIFO as it is generally the better algorithm. This is mainly due to the fact that the page accesses has shown little repetition.

(e)

In scenario 1, both CPU utilization and disk utilization is low, which indicates that the CPU is possibly bounded by IO bound processes and there is a lot of waiting time. CPU bound processes should be scheduled, if there is any, when the CPU is waiting for IO response.

In scenario 2, CPU utilization is low but the paging disk utilization is high. Thrashing occurs, and the operating system should consider allowing the current process to use more memory space.

In scenario 3, CPU utilization is high but the paging disk utilization is low. This is good, and no additional measures has to be taken.

Do not write below this line

Statement of good academic conduct

By submitting this assignment, I understand that I am agreeing to the following statement of good academic conduct:

- I confirm that this assignment is **my own work** and I have not worked with others in preparing this assignment.
- I confirm this assignment was written by me and is in my own words, except for any materials from published or other sources which are clearly indicated and acknowledged as such by appropriate referencing.
- I confirm that this work is not copied from any other person's work (published or unpublished), web site, book or other source, and has not previously been submitted for assessment either at the University of Birmingham or elsewhere.
- I confirm that I have not asked, or paid, others to prepare any part of this work for me.
- I confirm that I have read and understood the University regulations on plagiarism (<https://intranet.birmingham.ac.uk/as/registry/policy/conduct/plagiarism/index.aspx>).