

Revision week

Neural NLP – from embeddings to LLMs

Venelin Kovatchev

Lecturer in Computer Science

v.o.kovatchev@bham.ac.uk

Outline

- Word embeddings
- Compositionality and end-to-end networks
- Encoder-decoder and attention
- Transformers
- Transfer learning and types of transformers
- LLMs and RLHF finetuning

Vector representations and word embeddings

The distributional hypothesis and word embeddings

- The distributional hypothesis in semantics

- What does it state?

distributional

- How does it affect NLP?

- Word embeddings

- Encoding words as "semantic" vectors

Count based vector representations

- Obtain a large corpus in the language/domain of interest
- Define a context of co-occurrence
 - What contexts can you think of?
- Count and fill in a co-occurrence matrix
- Apply transformations (which?)

tf-idf, dimensionality reduction

Word2Vec

- Learning embeddings directly from text


- A simple neural architecture

- Two algorithms

- What is the difference between them?


predicting a word

CBow



pick a word, predict a surrounding word.

Skip-gram



Negative Sampling

- Global objective: maximize the log probability of the dataset of size T with a context size c

$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

- Calculate the probability of every word given context (or the other way around)

$$p(w_O | w_I) = \frac{\exp(v'_{w_O}{}^T v_{w_I})}{\sum_{w=1}^W \exp(v'_w{}^T v_{w_I})}$$

- Training with a softmax over the whole vocabulary is expensive
- Convert the task into “classifying the correct objective” using a logistic

$$P(+|w, c) = \sigma(\mathbf{c} \cdot \mathbf{w}) = \frac{1}{1 + \exp(-\mathbf{c} \cdot \mathbf{w})}$$

Calculating similarity

- How do we calculate similarity between vectors?
- Dot product

$$\text{dot product}(\mathbf{v}, \mathbf{w}) = \mathbf{v} \cdot \mathbf{w} = \sum_{i=1}^N v_i w_i = v_1 w_1 + v_2 w_2 + \dots + v_N w_N$$

- Cosine

$$\text{cosine}(\mathbf{v}, \mathbf{w}) = \frac{\mathbf{v} \cdot \mathbf{w}}{|\mathbf{v}| |\mathbf{w}|} = \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$

Post-training word embeddings

- We train embeddings using a logistic classification (with negative sampling)
- What happens with the logistic after the training?
- Can you think of another algorithm that uses similar approach?

ELMO

- What is the most important difference between ELMO and word2vec?

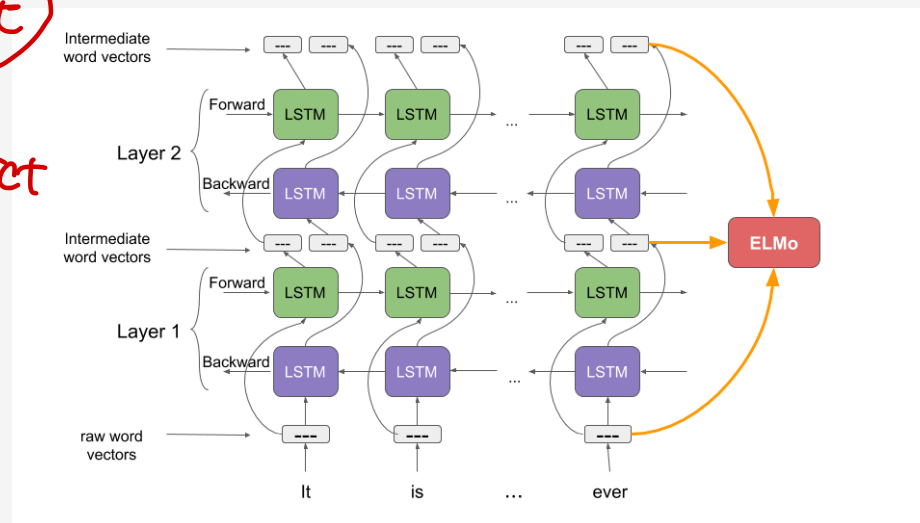
W2V static (same representation for the same word)

ELMO (different based on its context)

- What makes ELMO representations "deep"?

ELMO uses the whole network to predict

- What is the training objective behind ELMO?



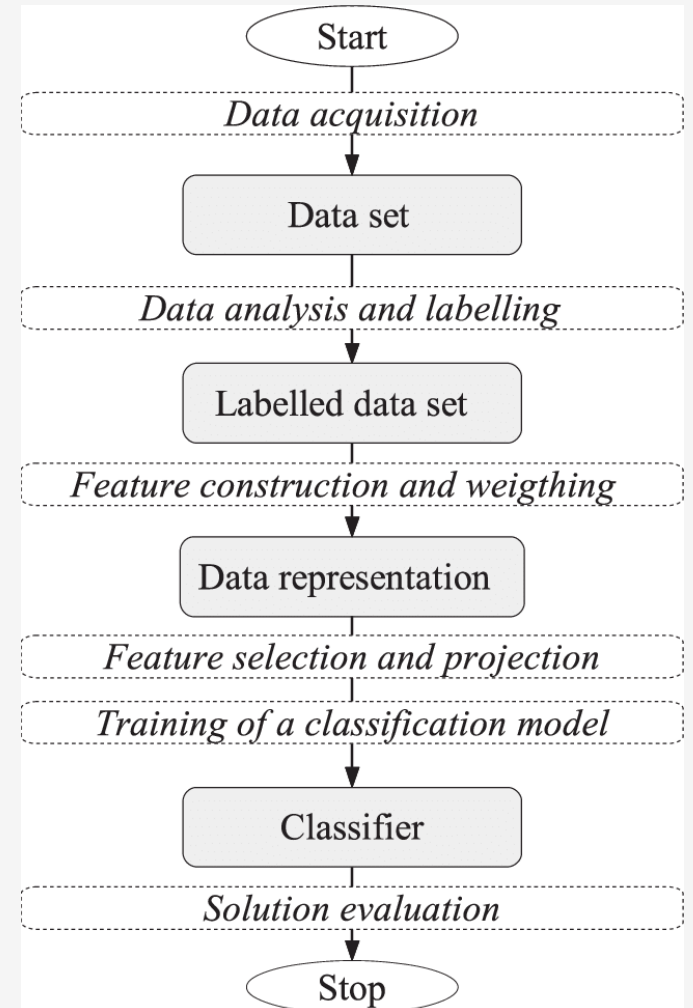
Compositionality and End-to-end

Feature engineering

- Analyze the problem, the input, and the desired outcome
- Explore existing resources and processing techniques
- Select the most relevant features and feature-extraction methods
- Empirically test what works best

Text classification using features

- Step by step process
- Involves active human engagement
- Feature selection and extraction
- Data is fed into a classifier (Logistic, NB, SVM)
- Iteratively improve feature selection and model (hyper) parameters



Why not go end to end?

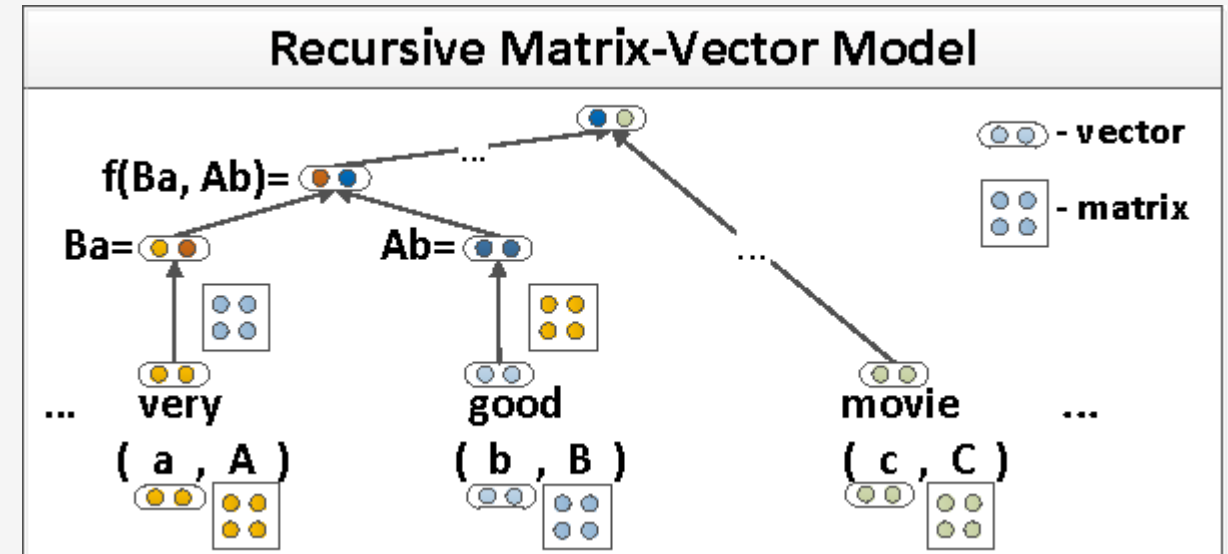
- Do we need full pipelines?
- Embeddings make it possible to “feed” text directly into models
- Is it possible to go fully end-to-end and eliminate
 - Accumulation of errors
 - Human labor and supervision
 - (In)compatibility issues between elements?

Embeddings and the problem of compositionality


- Embeddings represent individual words
- NLP deals with processing texts
- How to go from word representations to text representations?

Composing word meaning

- Vector operations
 - Vector addition
 - Pointwise vector multiplication
 - Vector concatenation
 - Complex matrix-vector operations
- Which of these operations consider text structure?
- Let a neural network do the compositionality




Multi-layer perceptron



- $y = \text{softmax}(h_1 \cdot W_2 + b_2)$
- $h_1 = f(h_0 \cdot W_1 + b_1)$
- $h_0 = f(x \cdot W_0 + b_0)$
- Non-linear functions f :
 - Sigmoid: $\sigma(x) = \frac{1}{\exp(-x)}$
 - Hyperbolic: $\tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)}$
 - ReLU: $\text{rect}(x) = \max(0, x)$


Recurrent neural networks

- How to represent text of varying length?
- Copy the network for each word
- At each timestep t , input is X_t and $h_{(t-1)}$
- I + like + tacos + "al pastor"
- Left to right, combining words one at a time
- Sequence classification
- Sequence to sequence



RNNs (formally)

- At each timestep
 - Input x_t and previous hidden state $h_{(t-1)}$
 - Three sets of weights:
 - input (W_x), hidden (W_h) , and output (W_o)
 - $h_t = f_h(W_x x_t + W_h h_{(t-1)})$
 - $O_t = f_o(W_o h_t)$
 - $J_t = f(O_t, y_t)$
- Pop-quiz: What are we predicting at O?



Neural language modeling with RNNs

- More formally:
- $h_t = \tanh(W_x x_t + W_h h_{(t-1)})$
- $\hat{y}_t = \text{softmax}(W_y h_t)$
- $J_t = -\log \hat{y}_{t, \text{correct}}$ (- log prob the word at t+1)

- $J_{\text{sent}} = \frac{1}{T} \sum_{t=1}^T -\log \hat{y}_{t, \text{correct}}$

- Pop quiz: What is "weight tying" in RNNs?

Sharing the weights in RNN
Same

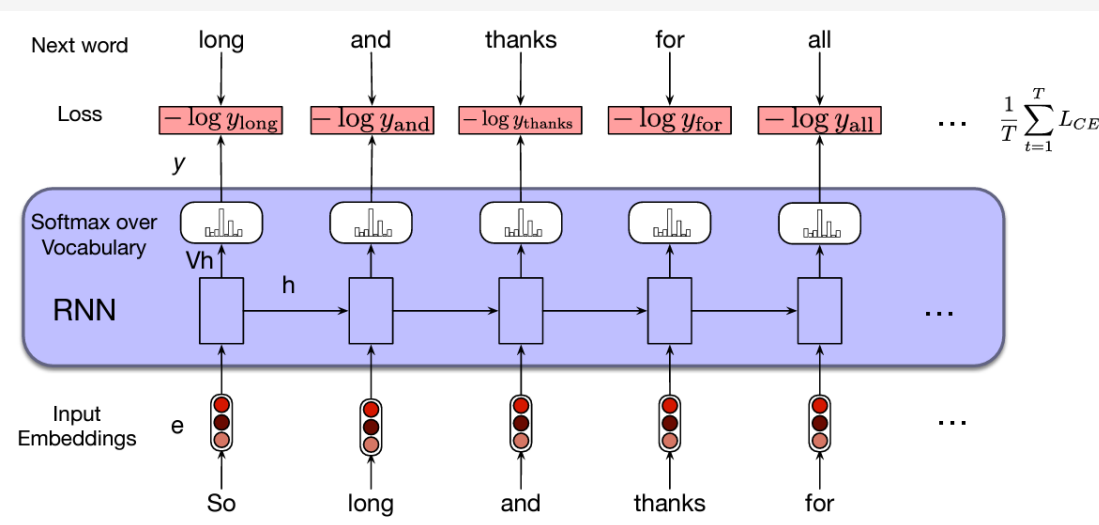


Figure 9.6 from SLP, chapter 9

Stacked and bidirectional RNNs


- The basic RNN is a powerful tool
- We can go deeper
- Stacking RNNs
- Bidirectional RNNs

LSTM

- Popular variation of RNN that address native limitations
- Same recurrent concept (copy the network)
- Three different "gates":
 - Forget gate
 - Input gate
 - Output gate
- Gates manipulate the flow of information through time
 - Addressing conflicting objectives

Understanding the gates


- All gates have the same format:
 - A feedforward layer followed by a sigmoid
- The gates are filtering out information at a certain part of the network
- Each gate has two important aspects:
 - How to calculate the filter
 - What is the input that the filter is applied to



Convolutional neural networks (CNN)

- Another popular architecture for NLP
 - Deals with text of various size
- Inspired by computer vision success
- Applying filters of different size to the input (2,3,4) + pooling
 - Analogous to n-grams

Convolutional neural networks (CNN)




Compositionality

- How each network “composes” meaning
 - Feed-forward network
 - Linear combination of words (no order) + activation function
 - RNN/LSTM
 - Recursively, word by word (linear order)
 - CNN
 - Locally, similar to n-gram (proximity window, limited order importance)

Encoder-decoder and attention

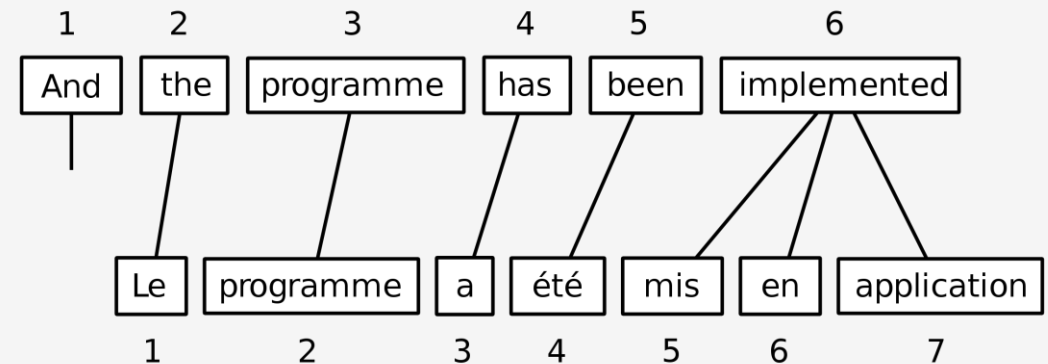
Sequence labeling vs sequence-to-sequence problems

- Consider the following two tasks



- What is similar between them?

- What is different?



- How would you approach each of them?

Key differences

- Same length vs different length
- One-to-one alignment vs no one-to-one alignment
- Local dependencies vs long-distance dependencies
 - Within the output
 - Between the input and the output

Encoder decoder

- We use a model family called encoder-decoder
- Simple idea
 - Encoder “represents” the source (e.g., English)
 - Decoder “generates” the target (e.g., German)
- Can you suggest tasks that can use encoder-decoder?





Fig 9.16

Using separate RNNs for encoder and decoder

- Train two models
- Pass the context at every step


$$\mathbf{h}_t^d = g(\hat{y}_{t-1}, \mathbf{h}_{t-1}^d, \mathbf{c})$$

- Can you point a potential problem?
- What could improve this architecture?
- What is the purpose of the encoder?
- Should it be able to generate?



Encoder decoder across modalities: image captioning


- The encoder and decoder “talk” via the context
- They don’t have to be the same type of model
- The modalities don’t have to match
 - Speech to text
 - Image to text



Deep Visual-Semantic Alignments for Generating Image Descriptions

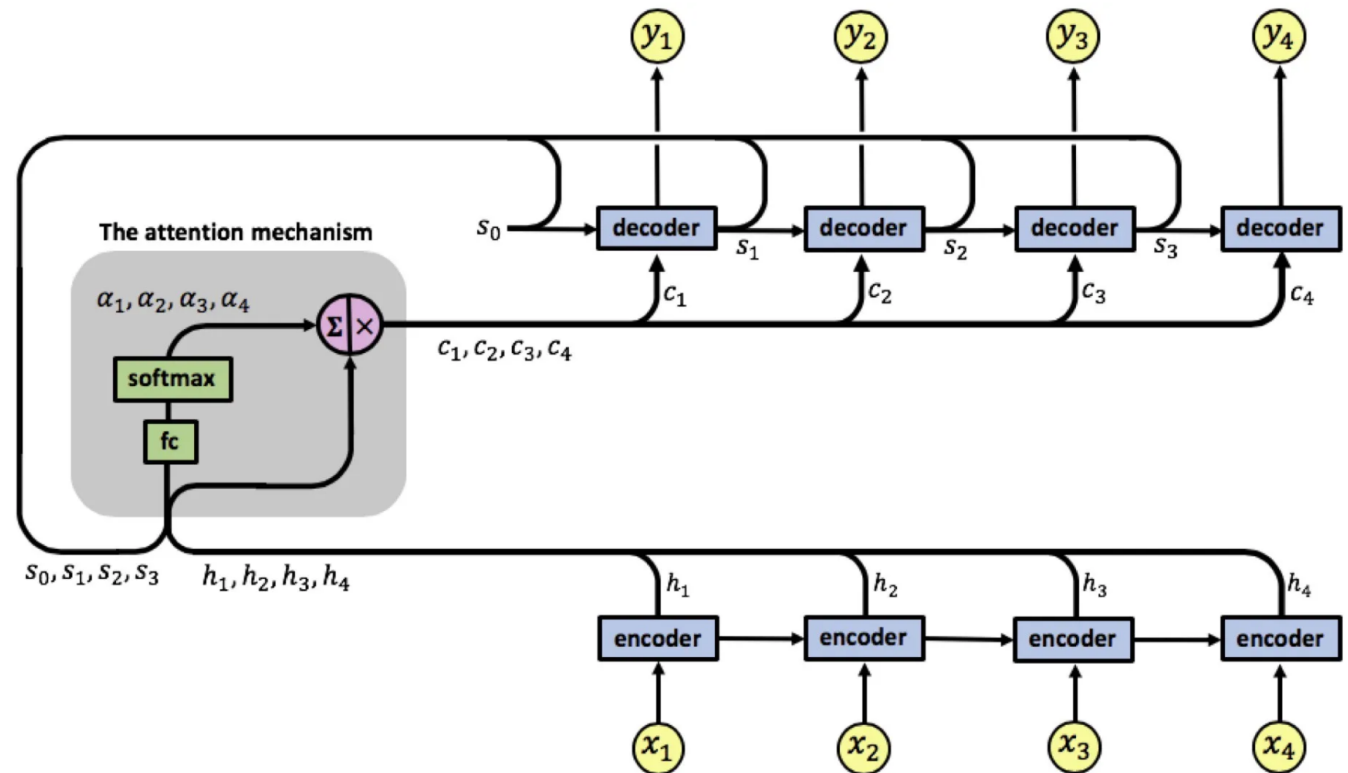
Attention – basic implementation

- Intuition: each token in the target should use a “personalized” context
 - Should have access to all hidden states in the encoder
 - The context should have a fixed (vector) length
- Weighted sum of all encoder hidden states
 - Calculated separately at each decoder step
 - Using the hidden state at (t-1)
- Dot product attention
 - Calculate the similarity between $h_{(t-1)}$ and each encoder state h^e
 - Use the similarity scores to calculate the weighted sum




Visualization of RNN with attention

- RNN with attention
- Attention is learned via a simple FFN



Visualizing attention


- Linear weights are interpretable
- We can see which word is more important
- Can we use attention for explainability?



Transformers


Elmo architecture

- How can we improve over that?



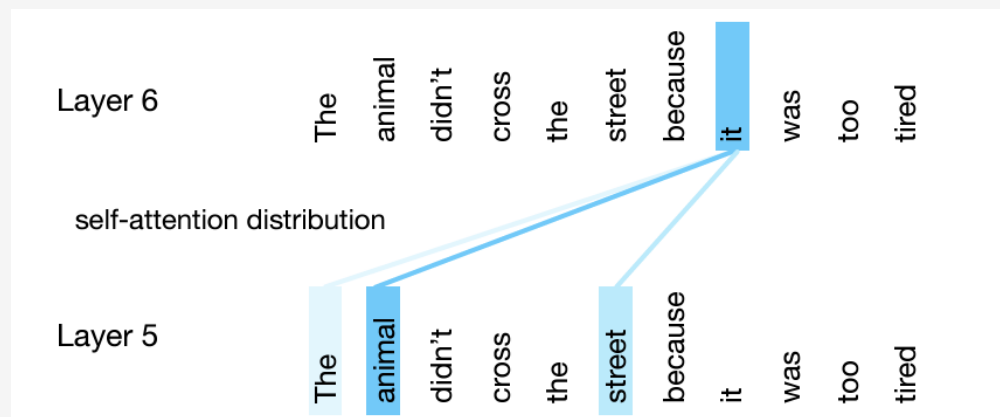
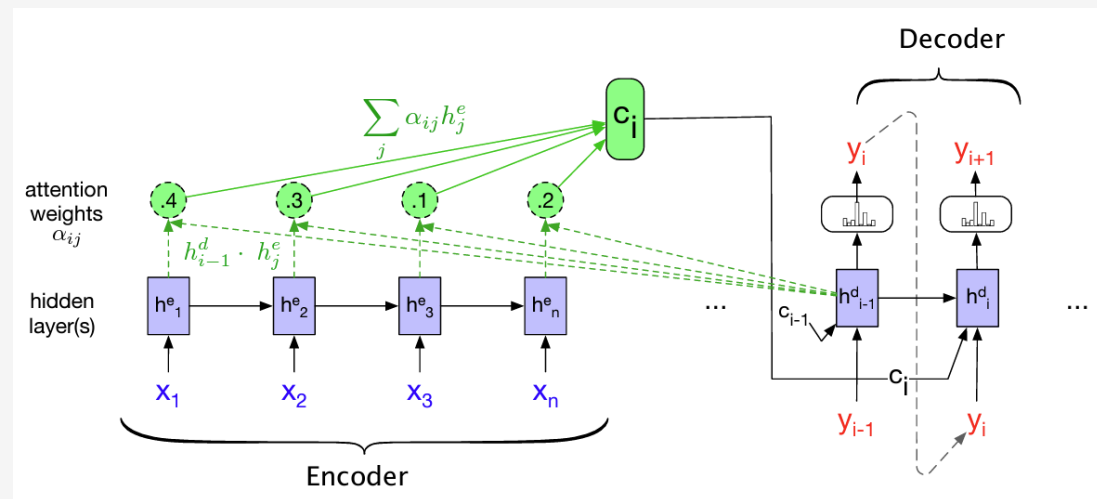
Self attention

- Attention works better than RNN/LSTM for encoder-decoder models
- Can we use attention for a standalone network?



Self attention (2)

- Self attention is a key concept in building transformers
- It applies the same approach as attention in encoder-decoder, but on itself



Causal self attention (intuition)


- Similar to RNNs, we have a 1:1 input-output mapping
- Same basic approach as original attention
- Dot product + softmax + weighted sum

$$\text{score}(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i \cdot \mathbf{x}_j$$

$$\begin{aligned}\alpha_{ij} &= \text{softmax}(\text{score}(\mathbf{x}_i, \mathbf{x}_j)) \quad \forall j \leq i \\ &= \frac{\exp(\text{score}(\mathbf{x}_i, \mathbf{x}_j))}{\sum_{k=1}^i \exp(\text{score}(\mathbf{x}_i, \mathbf{x}_k))} \quad \forall j \leq i\end{aligned}$$

$$\mathbf{a}_i = \sum_{j \leq i} \alpha_{ij} \mathbf{x}_j$$

- Which is the most similar token to \mathbf{x}_3 ? What is the input to the first hidden layer?



Decomposing input vectors

- We can use simple attention and it works
- Transformers introduce query, key, value
- What are they, why do we need them and how do we use them?
 - The “dictionary” analogy
 - A semantic explanation, grounded in NLP


How to model asymmetric compositionality in attention?

- Classical attention (that we have seen) has 1:1 correspondence

- Dot product attention is commutative

- $a \cdot b = b \cdot a$

- $\text{score}(\text{"black"}, \text{"dog"}) = \text{score}(\text{"dog"}, \text{"black"})$



- Pop quiz: would "black" have the same importance on "dog" as "dog" would have on "black"?

The query, key, value

- We project the input vector x to three vectors that serve different purpose: "query", "key", and "value"
- Two vector operations in the original attention:
 - "Score": for indexes i and j , calculate how important is x_j for x_i : $\text{score}(x_i, x_j)$
 - "Scale": for index i , calculate the hidden state h_i as a weighted sum of $x_1 \dots x_i$: $h_i = \sum_{j \leq i} \alpha_{ij} x_j$
- Each input vector x can have three different roles
 - Argument 1 in $\text{score}()$ ["dog" in $\text{score}(\text{"dog"}, \text{"black"})$] -> **query**
 - Argument 2 in $\text{score}()$ ["black" in $\text{score}(\text{"black"}, \text{"dog"})$] -> **key**
 - The **value** used in scale to calculate the hidden state

Query, Key, Value (formally)

- We learn three different matrices (W^Q, W^K, W^V)
- Every input vector x_i is projected to three different representations
 - $q_i = x_i W^Q ; k_i = x_i W^K ; v_i = x_i W^V$
- The new formula for score: $\text{score}(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{q}_i \cdot \mathbf{k}_j$
- The new formula for calculating weights: $\mathbf{a}_i = \sum_{j \leq i} \alpha_{ij} \mathbf{v}_j$
- Pop quiz: which token will have the most impact on x_3 ?


The transformer self attention

1. $\mathbf{q}_i = \mathbf{x}_i \mathbf{W}^Q; \mathbf{k}_i = \mathbf{x}_i \mathbf{W}^K; \mathbf{v}_i = \mathbf{x}_i \mathbf{W}^V$

2. and 3. $\text{score}(\mathbf{x}_i, \mathbf{x}_j) = \frac{\mathbf{q}_i \cdot \mathbf{k}_j}{\sqrt{d_k}}$

4. $\alpha_{ij} = \text{softmax}(\text{score}(\mathbf{x}_i, \mathbf{x}_j)) \quad \forall j \leq i$

5. and 6. $\mathbf{a}_i = \sum_{j \leq i} \alpha_{ij} \mathbf{v}_j$



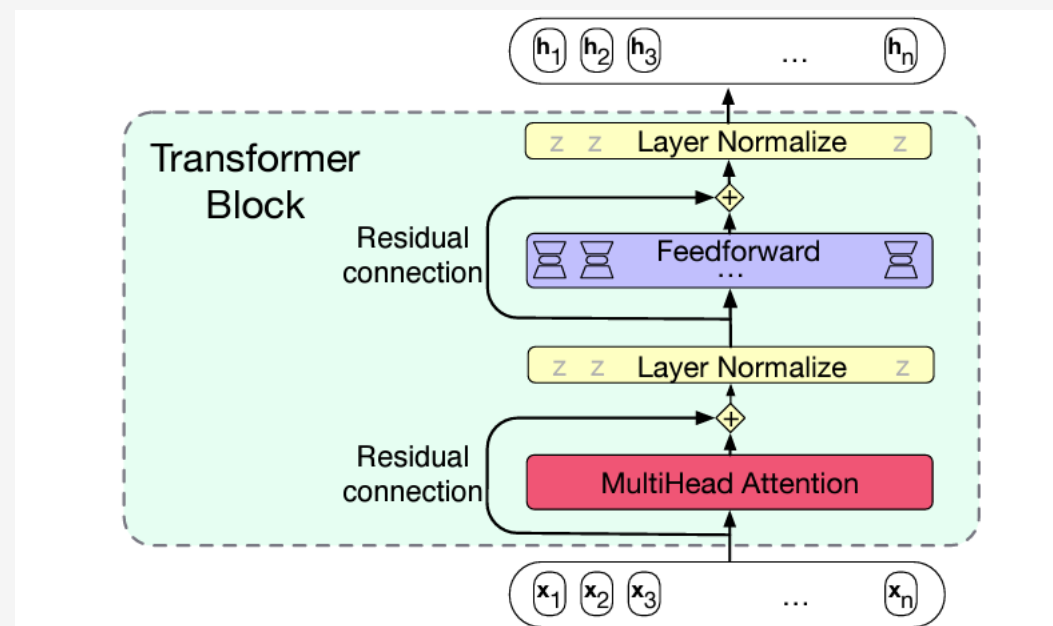
Multiheaded self-attention

- Instead of using a single self attention, we can use multiple
 - Each "head" has its own weights W^Q, W^K, W^V
 - The outputs of all heads are concatenated and projected to input dimensions
 - You can also think of multiheaded attention as "breaking" one big attention into specialized subsets
- Formally:

$$\begin{aligned}\mathbf{Q} &= \mathbf{XW}_i^Q ; \mathbf{K} = \mathbf{XW}_i^K ; \mathbf{V} = \mathbf{XW}_i^V \\ \mathbf{head}_i &= \text{SelfAttention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) \\ \mathbf{A} &= \text{MultiHeadAttention}(\mathbf{X}) = (\mathbf{head}_1 \oplus \mathbf{head}_2 \dots \oplus \mathbf{head}_h) \mathbf{W}^O\end{aligned}$$

The transformer block

- Residual connection
 - Copy the input of a layer to its output
- Layer normalize
 - Rescale each x vector to 0-mean with $STD=1$
- Positional feedforward
 - Apply the same fully connected FFN to each x




The transformer block (formally)

- Simplified representation
 - $O = \text{LayerNorm}(X + \text{MultiHeadAttention}(X))$
 - $H = \text{LayerNorm}(O + \text{FFN}(O))$
- You can change the order of operations in some implementations

Encoding the Input. Positional Embeddings.


- Semantic embeddings
 - One-hot encoding maps to a row in a matrix
- Positional embeddings
 - One embedding for each position
 - Learnable; Same dimension as semantic
- Add semantic and positional embeddings



- Alternative techniques: use functions (sine/cosine); calculate relative positional embeddings

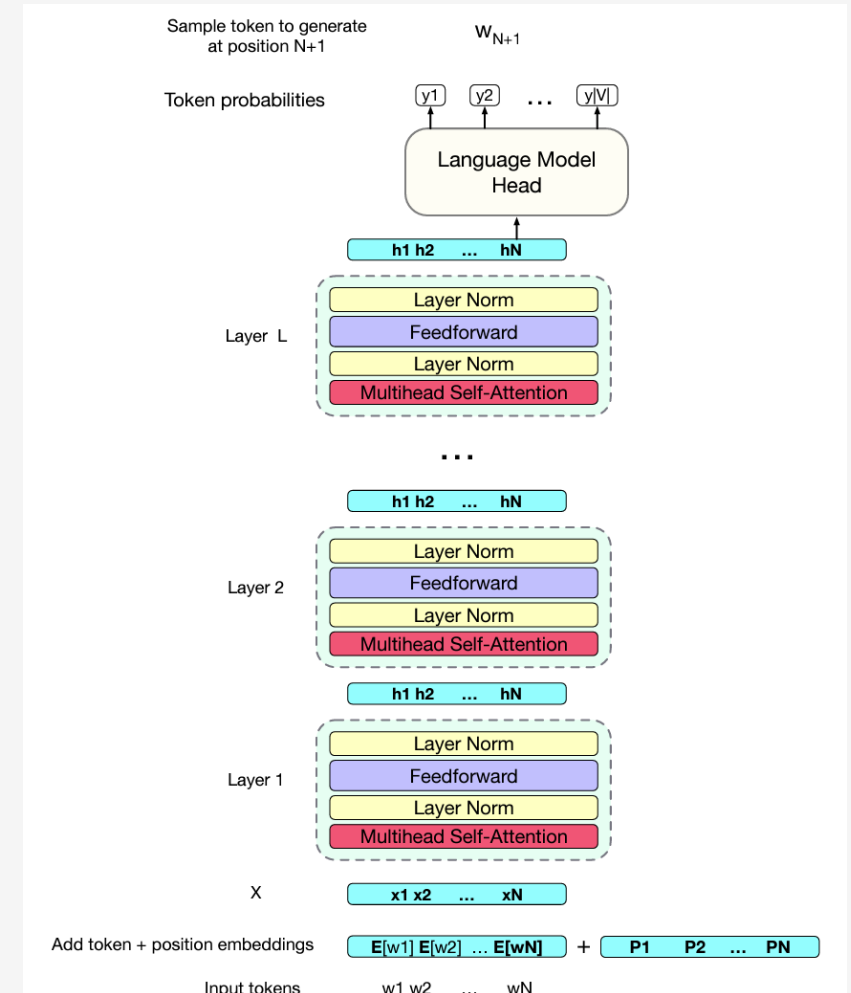
Language modeling head

- Language modeling
 - Efficient for learning representations
 - Self-supervised
- Project h_N to vocabulary size
 - Do we know any computational tricks for that?
 - What would h_N^L look like?



A final transformer representation for LM

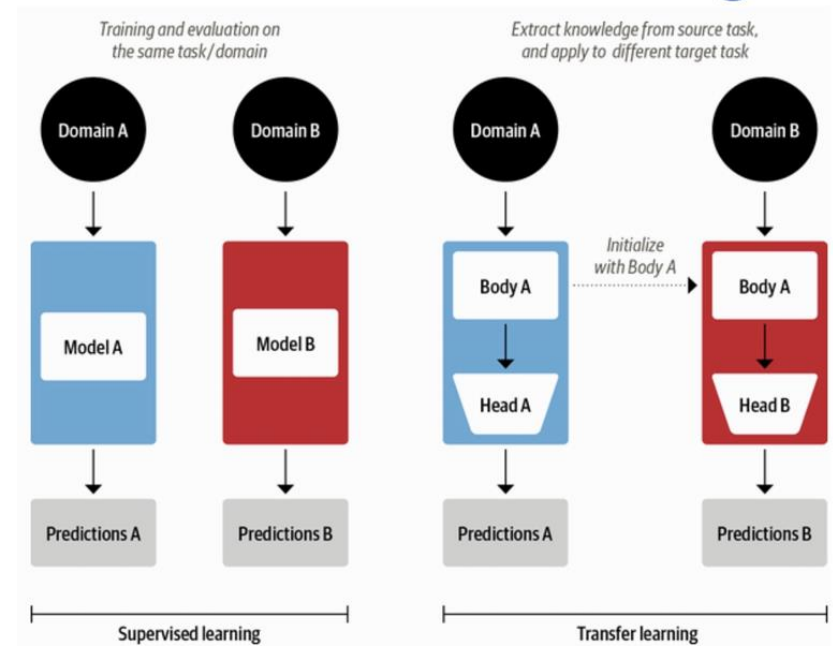
- Token + positional embedding
- Multiple stacked transformer blocks
- A classification head
- Language modeling with weight tying and sampling



Transfer Learning and Types of Transformers

Supervised learning vs Transfer learning

- What are the goals and benefits of transfer learning?
- What are some potential issues or risks?
- Are there any other paradigms that you can think of in that area?



Source: Lewis Tunstall, Leandro von Werra, and Thomas Wolf (2022), Natural Language Processing with Transformers: Building Language Applications with Hugging Face, O'Reilly Media.

The decoder transformer: GPT

- GPT1 combines different concepts we know so far
 - The standard transformer block
 - Neural Language Modeling
 - Transfer learning capabilities
- Pop quiz: What kind of attention does it use?
- Intuition:
 - Generative pre-training
 - Discriminative finetuning

Finetuning GPT

- After pretraining, use the hidden state at last layer
- Add a last linear layer with m neurons (m = number of classes)

- Predict the target class:


$$P(y|x^1, \dots, x^m) = \text{softmax}(h_l^m W_y).$$

- Maximize the probability of the correct labels (need labeled data)

$$L_2(\mathcal{C}) = \sum_{(x,y)} \log P(y|x^1, \dots, x^m).$$

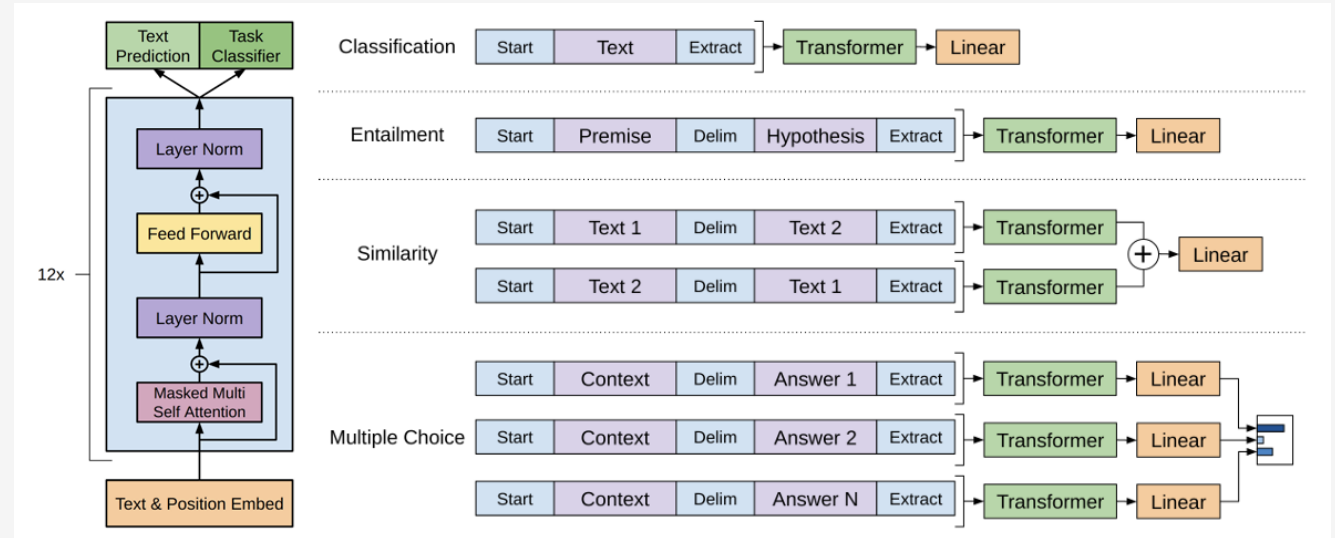
- Combining both losses together

$$L_3(\mathcal{C}) = L_2(\mathcal{C}) + \lambda * L_1(\mathcal{C})$$



Task specific input transformations

- Task reformulating
- Using special tokens (sep, start/end)
- Comparing separate "streams"
- Task design is a non-trivial task
 - Task formulation; Data format; Metrics and Evaluation




The encoder transformer: BERT (main difference is type of attention)

- The original encoder-only transformer
- An English-only sub-word vocabulary consisting of 30,000 tokens
 - Most of the modern algorithms use subwords tokenizers and embeddings
- 768 hidden size
- 12 layers, 12 heads in each multi-head attention
- 100M parameters
- Trained on two tasks: Masked Language Modeling and Next Sentence Prediction

The encoder transformer

- The encoder in “attention is all you need”
- Same architecture as the decoder
- Bi-directional self-attention
 - All key/query values, no masking
- Better for encoding source information



Masked language modeling objective

- Based on “cloze” tasks:
 - “Can I have a ____ of water, please?”
 - Does that remind you of something?
- Masked Language Modeling (MLM)
 - Randomly sample tokens from the text and perform alternations
 - Predict the original inputs for each position

Next Sentence Prediction



- MLM predicts relationships between words
- Transformers want to also process sentences
- Next sentence prediction task
 - Given two sentences, predict whether they are a pair of adjacent sentences

Next sentence prediction. The CLS token.

- Next sentence prediction
 - 50 % true adjacent pairs
 - A special [CLS] token added at the beginning
 - A special [SEP] token added between texts
 - Special "sentence position" (first/second) are added to input
- When predicting the sentence relation, we use the CLS as an input to softmax

Adapting other tasks to work with BERT

- How would we perform paraphrase identification?
- What is the input/output/classification process?
- Performing other tasks:
 - Extractive QA
 - Sequence labeling



The encoder-decoder model

- Encoder
 - Bi-directional attention can “see” all tokens
 - Follows the architecture we have seen last week
- Decoder
 - Causal attention
 - Additional Multiheaded Attention
 - Why do we need it?
 - What would be the Q, K, V used by it?

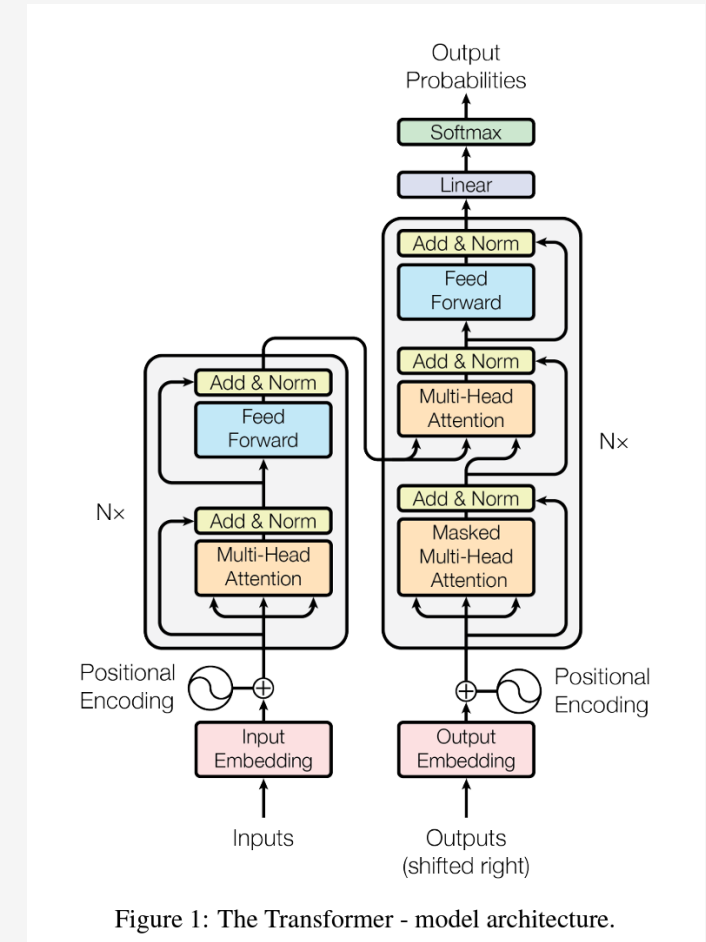
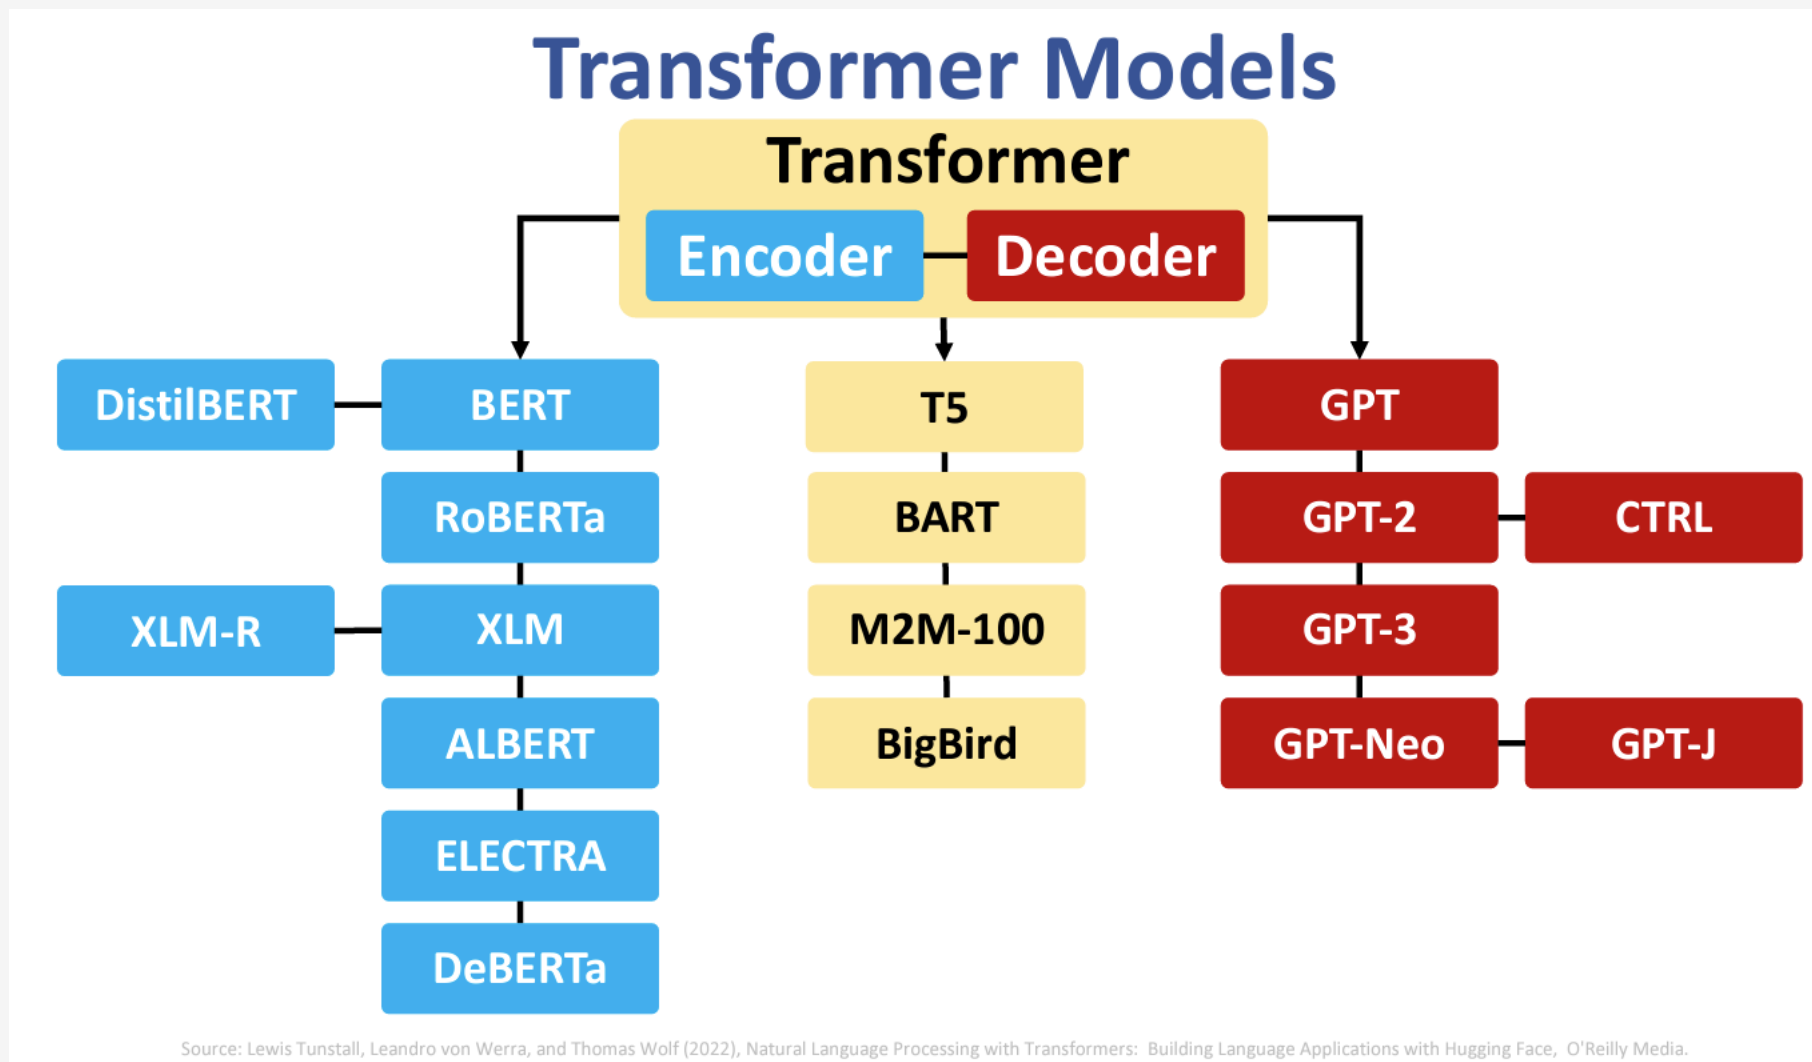


Figure 1: The Transformer - model architecture.

The transformer family tree



In-context learning and RLHF finetuning

In-context learning

- Taking transfer learning to the extreme
- Using the input to specify the task
 - "What is the sentiment of the following text: I like this movie, it's the best in the Avengers series!"
 - "Do those sentences contradict each other: I bike to work every day. <SEP> I drive to work every day."
- Emerging property
 - A by-product of scaling the model above a certain size

Zero- One- and Few-shot learning

- Three different experimental conditions
- No gradient update or finetuning
- The only difference – number of examples

The three settings we explore for in-context learning

Zero-shot

The model predicts the answer given only a natural language description of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 cheese => ..... ← prompt
```

One-shot

In addition to the task description, the model sees a single example of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 sea otter => loutre de mer ← example
3 cheese => ..... ← prompt
```

Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 sea otter => loutre de mer ← examples
3 peppermint => menthe poivrée ←
4 plush girafe => girafe peluche ←
5 cheese => ..... ← prompt
```

InstructGPT – training models to follow instructions

- Scale is not everything
 - Hallucinations
 - Toxicity
 - Lack of helpfulness
- Improve the training (and evaluation) procedures
- “Align” models with their users

The LM training objective

- Language modeling is not “following instructions”
- Language modeling does not take (individual) preferences
- Every training sequence is equally important
- Preference in output (in language modeling) depends on
 - The observed frequency in training data
 - The sampling strategy
- Few- and Zero-shot learning are an “emerging” side effect, not an intentionally defined goal

The training process of InstructGPT


Step 1

**Collect demonstration data,
and train a supervised policy.**


A prompt is
sampled from our
prompt dataset.



A labeler
demonstrates the
desired output
behavior.




This data is used
to fine-tune GPT-3
with supervised
learning.




Step 2

**Collect comparison data,
and train a reward model.**


A prompt and
several model
outputs are
sampled.



A labeler ranks
the outputs from
best to worst.



This data is used
to train our
reward model.



Step 3

**Optimize a policy against
the reward model using
reinforcement learning.**


A new prompt
is sampled from
the dataset.



The policy
generates
an output.



The reward model
calculates a
reward for
the output.



The reward is
used to update
the policy
using PPO.

